

What Do Bug Repositories Tell Us About Software Development Workflow? An Exploratory Study

Rodrigo Souza and Christina Chavez
Software Engineering Labs
Department of Computer Science - IM
Federal University of Bahia (UFBA), Brazil
{rodrigo,flach}@dcc.ufba.br

Abstract—Background. Bug repositories have been increasingly studied with the objective of unveiling relevant knowledge about development practices and coordination tasks in software projects. Although there is plenty of publicly available repositories, researchers have raised issues regarding the accuracy of the data recorded in such repositories.

Aims. In this paper, we investigate some bug repositories in depth in order to understand the possible meanings behind the data (e.g., what does it mean to say that a bug has been “verified”?) for each particular software project analyzed.

Method. We have analyzed bugs from two large software projects, Eclipse and Netbeans. We have performed in-depth qualitative analysis within a small sample of bugs from both projects. Also, we have performed a broad textual analysis by applying latent Dirichlet allocation on all bug reports. Finally, we have used descriptive statistics in order to account for temporal (pre-release, post-release) and social (role of user: committer, fixer, reporter, end user) factors in the bug fixing workflow.

Results. We have found that statuses mean distinct things in different projects. Also, the kind of activity associated with a certain status change is dependent on the roles of the participants and the stage of the software release cycle. The results reinforce the threats involved in mining bug repositories and provide a way to estimate the inaccuracy in data. We recommend that future researchers use our method prior to mining bug repositories in order to select projects in which the meaning of the data records match their expectations.

Keywords—software evolution; software development practices; bug repositories; empirical study;

I. INTRODUCTION

Software development is a complex activity that demands the coordination of many people, such as programmers, designers, and managers. Software development activities should be coordinated in order to deliver products with high quality and on schedule. Many practices and processes have been proposed in order to aid the software development, targeting at high quality, low cost, fast delivery, or a combination of these.

Several researchers are concerned with the validity of studies based on bug tracking systems [1], [2], [3], [4]. The main concern is the accuracy of the information that is reported in such systems.

II. BACKGROUND

A. Bug Tracking Systems

Bug workflow. Bugzilla. [2]

Repositories of bugs. Arquivam relatos de defeitos e solicitações de mudança. Tais repositórios ajudam os desenvolvedores a registrar, discutir e acompanhar a situação de cada solicitação, desde a atribuição de uma solicitação a um desenvolvedor até a verificação da correção ou melhoria. Exemplos incluem o Bugzilla, o Trac e o Jira.

Bug tracking systems are applications that allow users and developers of a software system to manage a list of bugs for the system. Usually, users and developers can report bugs, along with information such as steps to reproduce the bug and the operating system that was used. Then, developers choose bugs to fix and can report on the progress of the bug fixing activities, ask for clarification etc.

One important feature of a bug that is recorded on bug tracking systems is its status. The status records the progress of the bug fixing activity, and, as such, provides data to software engineering studies. Figure 1 shows each status that can be recorded in Bugzilla, a popular bug tracking system, along with typical transitions between statuses, i.e., the workflow. Other systems present comparable status and transitions.

In the simplest cases, a bug is created and receives the status UNCONFIRMED or NEW. Next, it is assigned to a developer, and then it is resolved, possibly by fixing it with a patch on the source code. The solution is then verified by someone in the quality assurance team. If it passes the quality requirements, it is closed. If it doesn't pass the quality requirements or if the solution only partially fixes the problem, the bug is reopened.

B. The Eclipse Project

History.

Recommendations about Bugzilla usage. Four-eyes principle.

The Eclipse Project was created in 2001 by IBM and released into open source in 2004¹, with focus on the

¹<http://www.eclipse.org/org/>

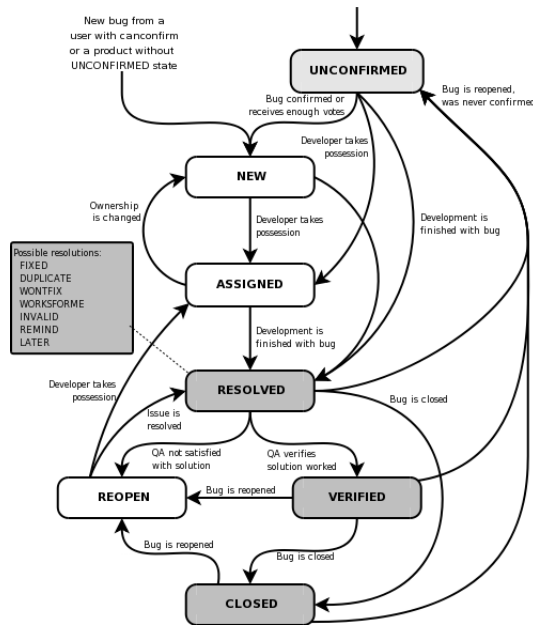


Figure 1. Workflow for Bugzilla. Source: <http://www.bugzilla.org/docs/2.18/html/lifecycle.html>.

creation of an extensible platform that supports the software development lifecycle. Today, it is maintained by its open source community, as well as by IBM employees and people from other organizations.

The Eclipse Project consists of many subprojects, including the Platform project and dozens of plugins. Although each subproject is maintained by a distinct team and is relatively autonomous, all of them should follow a general, documented process.

All subprojects use Bugzilla for bug reporting, and there are clear rules regarding the process of change in the software². Here is an excerpt of a document in the Eclipse web site that explains the bug fixing process³ (emphasis added):

When a new bug is entered it begins life with a Status of either **Unconfirmed** for normal users or **New** for users with commit privileges. The bug is typically assigned to the component owner. The component owner will usually use a query of Status = Unconfirmed or New and Assigned to = me to browse what is essentially the component's inbox. She or he will assign bug reports to developers. Email will be generated to the developer (this part is configurable). (Note that the act of reassigning a bug to somebody changes its status to New). The developer may change unconfirmed bugs to new.

The assigned developer will accept the bug which will change its status to **Assigned**. After working on the bug the developer will mark the bug as **Resolved** and will select a resolution (Fixed, Invalid, Wontfix, Worksforme).

Once a bug is resolved there are still a few states it can transition to. How you choose to use these states will be up to individual teams. In our example, when a test pass comes up a component owner can build a query that searched for resolved bugs. After testing that the fix worked a resolved bug can be transitioned to **verified**, directly to **closed**, or in fact, **reopened**. By searching for bugs with a Status of verified and a resolution of fixed developers can come up with release notes. A verified/fixed bug can then be transitioned to closed. And yes, closed bugs can be reopened if need be. As an added bonus other bug status can be transitioned through verified to closed as well. This gives developers the opportunity to test worksforme claims.

There is also a recommendation to follow the four eyes principle in the verification process (emphasis added):

When a developer fixes the bug, the status is set to **RESOLVED - FIXED**, and it is assigned to another developer on the team to verify. **It is important that the verifier be a different person than the fixer** because the fixer is too close to the code and thus may not be as diligent at testing the corner cases.

Finally, there are additional guidelines regarding the workflow:

- All bugs should be verified before the next integration build.
- When a committer verifies a fix, the status is changed to **VERIFIED**.
- When the project does a major release, the **VERIFIED** bugs are changed to **CLOSED**.

C. The Netbeans Project

History.

Recommendations for developers.

III. EXPERIMENTAL SETUP

Data extraction

Data processing:

- filtering (bugs up to december 20xx)
- classification script

Sampling:

- random sampling (for coding)
- entire data set

Analysis:

- grounded theory (open coding)

²http://wiki.eclipse.org/Development_Resources/HOWTO/Bugzilla_Use

³http://wiki.eclipse.org/Bug_Reporting_FAQ

- latent Dirichlet allocation
- descriptive statistics

We have had access to Eclipse Bugzilla database from October, 2001 to June, 2010, in MySQL format. The data was made available as part of the Mining Software Repositories 2011 Challenge⁴.

The database includes data from 155 Eclipse subprojects. However, we have discarded subprojects in which no bug was ever reopened, and subprojects with less than 50 reported bugs.

For the remaining 34 subprojects, we have selected only bugs that were verified until December 2009. The rationale is that if, within six months, a verified bug was not reopened, it probably will not be reopened after that and it is, thus, fully fixed.

IV. RESULTS

O que quer dizer a soluo de um bug foi verificada?

Como j foi comentado anteriormente, o Bugzilla permite que os relatios de bugs sejam marcados como RESOLVED (resolvido) e, depois, como VERIFIED (verificado). Mas afinal, o que significa verificar a soluo de um bug? Que atividades ou checagens foram realizadas antes de se marcar o bug como VERIFIED?

A documentao do Bugzilla tem algo a dizer sobre isso o status VERIFIED:

QA [quality assurance team] has looked at the bug and the resolution and agrees that the appropriate resolution has been taken.

Essa definio sugere que a verificao deve ser feita por uma equipe especializada em controle de qualidade. Alm disso, o verificador deve concordar que a resoluo apropriada. A definio do Bugzilla para por a, mas eu consigo imaginar alguns critrios para considerar uma resoluo apropriada:

O bug no se manifesta aps a aplicao da soluo. A soluo no gera novos bugs. O cdigo da soluo segue as convenes de codificao e polticas do projeto. A soluo possui uma boa relao custo-benefcio isto , ou no existem solues melhores ou a soluo um paliativo por conta da alta complexidade de uma soluo melhor. Minha expectativa, portanto, que, antes de o bug ser marcado como VERIFICADO, ele passe pelas seguintes atividades:

Aplicao do patch com a soluo do bug na verso do software onde o bug foi encontrado. Execuo do software com o bug, seguindo os passos para reproduo do bug, e verificando que o bug no se manifesta Inspeo do cdigo-fonte. Deve ser verificada a aderncia s convenes de codificao do projeto e a outras polticas (por exemplo, cada patch deve vir com testes de unidade). Execuo de testes de unidade (se houver). Deve se verificar que todos os testes passam. Bom, mas isso a teoria. Quais atividades so executadas na prtica, em projetos reais, antes de se marcar um bug como VERIFIED? Para

responder a essa pergunta, analisei um total de 80 bugs, considerando os projetos Eclipse/EMF, Eclipse/Platform, Netbeans/versioncontrol e Netbeans/profiler. Eis alguns casos verificados nesses projetos:

O reporter (ou outro usurio que tambm enfrenta o problema) verificou que, em uma build do produto que inclui a soluo, o bug deixou de se manifestar. Aparentemente no h inspeo do cdigo-fonte, e nem uma equipe de QA. Situao frequente no Netbeans. Pode ser um procedimento necessrio quando os desenvolvedores no conseguem reproduzir o bug. A soluo est disponvel em um build que divulgado no site do projeto. Esse critrio usado pelo projeto Eclipse/EMF e foi decidido em reunio. O projeto usa o status FIXED para bugs cuja soluo est no CVS. Bugs que j foram resolvidos h muito tempo so marcados como verificados, por se entender que, se nenhuma problema ocorreu at ento, isso significa que a soluo adequada. A mesma pessoa que resolve o bug o marca como VERIFIED logo depois. Aparentemente no existe verificao de fato nesse caso. O desenvolvedor que resolve o bug solicita que outro desenvolvedor verifique a soluo. Comum no Eclipse/Platform. No bug 249436 do Eclipse/EMF, um patch submetido por um desenvolvedor, mas outro desenvolvedor acaba reescrevendo o patch, fornecendo evidncias de que o cdigo-fonte foi inspecionado. No bug 269789 do Eclipse/EMF, h evidncias da elaborao e da execuo de testes de unidade. A anlise sugere que o status VERIFIED significa coisas diferentes em diferentes projetos. A execuo de uma verso corrigida do software parece ser o mtodo predominante de verificao. Em alguns casos, o status VERIFIED no representa verificao, e sim a indicao de que a soluo est disponvel para o usurio final. Evidncias de inspeo de cdigo e execuo de testes de unidade s foram encontradas no projeto Eclipse/Platform.

V. CONCLUSION

ACKNOWLEDGMENT

The first author is supported by FAPESB under grant BOL0119/2010.

REFERENCES

- [1] J. Aranda and G. Venolia, "The secret life of bugs: Going past the errors and omissions in software repositories," in *Proceedings of the 31st International Conference on Software Engineering*, ser. ICSE '09. Washington, DC, USA: IEEE Computer Society, 2009, pp. 298–308. [Online]. Available: <http://dx.doi.org/10.1109/ICSE.2009.5070530>
- [2] C. Bird, A. Bachmann, E. Aune, J. Duffy, A. Bernstein, V. Filkov, and P. Devanbu, "Fair and balanced?: bias in bug-fix datasets," in *Proceedings of the the 7th joint meeting of the European software engineering conference and the ACM SIGSOFT symposium on The foundations of software engineering*, ser. ESEC/FSE '09. New York, NY, USA: ACM, 2009, pp. 121–130. [Online]. Available: <http://doi.acm.org/10.1145/1595696.1595716>

⁴<http://www.msrrconf.org/msr-challenge.html>

- [3] T. H. D. Nguyen, B. Adams, and A. E. Hassan, "A case study of bias in bug-fix datasets," in *WCRE*, G. Antoniol, M. Pinzger, and E. J. Chikofsky, Eds. IEEE Computer Society, 2010, pp. 259–268.
- [4] K. Ayari, P. Meshkinfam, G. Antoniol, and M. Di Penta, "Threats on building models from cvs and bugzilla repositories: the mozilla case study," in *Proceedings of the 2007 conference of the center for advanced studies on Collaborative research*, ser. CASCAN '07. New York, NY, USA: ACM, 2007, pp. 215–228. [Online]. Available: <http://doi.acm.org/10.1145/1321211.1321234>